

Parallel Programming Tutorial - OpenMP Basics

Amir Raoofy, M.Sc.

Chair for Computer Architecture and Parallel Systems (Prof. Schulz)

Technical University Munich


2. Mai 2018



TUM Uhrenturm

Solution for Assignment 2

Solution for Assignment 2



```

1 typedef struct {
2     void *image;
3     int chunk_size, max_iter;
4     int x_resolution, y_resolution;
5     double view_x0, view_x1, view_y0, view_y1;
6     double x_stepsize, y_stepsize;
7     int palette_shift;
8 } compute_args;

```

```

1 static int global_start;
2 #define CHUNK_SIZE 8
3 pthread_mutex_t mtx = PTHREAD_MUTEX_INITIALIZER;
4
5 void mandelbrot_draw( args ...) {
6     global_start = 0;
7     int i, t;
8
9     pthread_t *thread = (pthread_t*)malloc\
10         (num_threads * sizeof(*thread));
11     compute_args *args = (compute_args*)\
12         malloc(num_threads * sizeof(*args));
13
14     for (t=0; t < num_threads; t++) {
15         args[t].image= (void*) image;
16         args[t].chunk_size = CHUNK_SIZE;
17         // ... similar for other parameters ...
18         pthread_create(&thread[t], NULL, kernel, &args[t]);
19     }
20     for (t = 0; t < num_threads; t++)
21         pthread_join(thread[t], NULL);
22
23     free(thread); free(args);
24 }

```

Solution for Assignment 2 (Cont.)

```

1 void* kernell(void* arguments){
2     compute_args *args = (compute_args*) arguments;
3     int chunk_size = args->chunk_size;
4     unsigned char (*image)[x_resolution1][3]= (unsigned char (*)(x_resolution1)[3])args->image;
5     // ... same for the rest of arguments ...
6     int start; //local variable
7     for (;;) { //infinite loop
8         pthread_mutex_lock(&mtx);
9         if ( y_resolution - global_start < 1 ) { // if every row is processed unlock and come out
10             pthread_mutex_unlock(&mtx); break;
11         }
12         start = global_start; global_start += chunk_size; // set the start and increase global variable
13         pthread_mutex_unlock(&mtx);
14         if ( y_resolution - start < chunk_size ) // for the thread that works on the last chunk
15             chunk_size = y_resolution - start;
16         for (int i = start; i < start + chunk_size; i++)
17         {
18             for (int j = 0; j < x_resolution1; j++)
19             {
20                 // ... calculation of pixels ...
21             }
22         }
23     }
24 }

```

Hints for Assignment 3

Hints for Assignment 3

- Use a profiler! (see last session)
- Try to reduce the critical region. **That is the bottleneck!**
- Use `std::ref()` to pass arguments by reference to a task function
- Use the launch policy `std::launch::async` when using `std::async` to explicitly spawn new threads

OpenMP

Introduction to OpenMP

- OpenMP is an API for explicit shared-memory parallelism
- Supported by most compilers (gcc, icc, msvc, clang)
- Utilizes OS threading capabilities (e.g. Pthreads)
- Fully documented in the specification (see <http://www.openmp.org/mp-documents/OpenMP4.0.0.pdf>)
- Comprised of three programming layer components

1. Compiler Directives

- Spawning parallel regions
- Distributing loop iterations across threads
- Synchronization
- ...

2. Runtime Library Routines

- Setting/Querying the number of current threads
- Querying thread-id's and wall-clock time
- ...

3. Environment Variables

- Setting number of threads
- Binding threads to processors
- ...

Directives

Format

```
#pragma omp <directive name> <{clause, ...}>
```

- `#pragma omp`
Required for all OpenMP C/C++ directives
- `directive name`
A valid OpenMP directive
- `{clause, ...}`
Optional. Clauses can be in any order
- Most OpenMP constructs apply to a structured block


Example

```
#pragma omp parallel default(shared) private(i)
```

Parallel Region

```
#pragma omp parallel <{clause, ...}>
```

- A block of code that will be executed by multiple threads
- Number of threads defined by `#cpu`, clauses, or env. variables
- The reaching thread (0) creates a team of N threads (1, ..., $N-1$)
- At the end of a block there is an implicit join (barrier)
- There may be nested parallel regions



Parallel Region - Example

- `omp_get_thread_num()` returns the current thread number
- `omp_get_num_threads()` returns the number of threads

```
1  int main(int argc, char** argv) {  
2  
3      #pragma omp parallel  
4      {  
5          printf("Hello World from thread %d\n", omp_get_thread_num());  
6  
7          // only executed by main thread  
8          if (omp_get_thread_num() == 0)  
9              printf("Number of threads is %d\n", omp_get_num_threads());  
10     }  
11     return 0;  
12 }
```

./hello_world

Hello World from thread 1

Hello World from thread 0

Number of threads is 3

Hello World from thread 2

Parallel Region - Clauses

- `if (<scalar expression>)`
only executed multithreaded if scalar expr. evaluates to non-zero
- `private (<list>)`
each thread gets a copy of variables in a comma separated list (variables might be uninitialized)
- `firstprivate/lastprivate (<list>)`
same as `private`, but value is copied at the entry/exit
- `shared (<list>)`
variables in list are shared (no elements of structs or arrays)
- `default (shared | none)`
sets the default behaviour (none means that data sharing needs to be explicit)
- `reduction (<operator: list>)`
reduction operation and associated operand
- `num_threads (<integer expression>)`
sets the number of threads for the parallel region
- ...

for Directive

```
#pragma omp for <{clause, ...}>
```

- Worksharing construct to execute the immediately following loop by a team of threads
- Assumes that a parallel region has already been initiated
- There is an implicit barrier at the end of the loop
- Clauses:
 - `schedule (static|dynamic|guided|runtime|auto)`
sets the scheduling behaviour (see next slide)
 - `nowait`
threads do not synchronize after the parallel loop
 - `ordered`
iterations must have the same order as in a serial program
 - `collapse`
specifies the number of (nested) loops that shall be collapsed into a larger iteration space
 - `private, firstprivate...`

schedule clause

- `schedule (static, chunk_size)`

The iterations are divided into chunks of size *chunk_size* and assigned to the threads in round-robin fashion. When no chunk size is specified, the iterations are equally divided (at most one iteration per thread).

- `schedule (dynamic, chunk_size)`

The iterations are distributed to threads in chunks as the executing threads request them. Each thread executes a chunk of iterations, then requests another chunk, until no chunks remain. Each chunk contains *chunk_size* iterations, except for the last chunk. If no chunk size is specified, it defaults to 1.

- `schdule (guided, chunk_size)`

Similar to `dynamic`, but...

At the beginning the size of each chunk is proportional to the number of unassigned iterations divided by the number of threads, decreasing to 1 (or *chunk_size*, if specified).

- `schedule (auto)`

The scheduling decision is given to the compiler/runtime system.

- `schedule (runtime)`

The scheduling decision is deferred until run time, the `schedule` and `chunk_size` are taken from internal control variables.

for Directive - Example

```
1 #include <omp.h>
2 #include <stdio.h>
3 #include <string.h>
4 #include <stdlib.h>
5 #include <unistd.h>
6
7 const char *colored_digit[] = {
8     "\e[1;30;1m0", "\e[1;31;1m1", "\e[1;32;1m2", "\e[1;33;1m3", "\e[1;34;1m4", "\e[1;35;1m5", "\e[1;36;1m6", "\e[1;
9 };
10
11 int main(int argc, char** argv) {
12     unsigned int x_size = 80;
13     unsigned int y_size = 40;
14     unsigned long str_len = strlen (colored_digit [0]);
15     char *string_2D = (char*)malloc(x_size * y_size * str_len + y_size);
16
17     #pragma omp parallel for schedule(runtime)
18     for (unsigned long i = 0; i < y_size; i++) {
19         for (unsigned int j = 0; j < x_size; j++) {
20             memcpy(string_2D + ( i * x_size * str_len + i ) + (j * str_len), colored_digit[omp_get_thread_num()], str_
21         }
22     }
```

```
OMP_NUM_THREADS=4 OMP_SCHEDULE="STATIC" ./scheduling
```

```
OMP_NUM_THREADS=4 OMP_SCHEDULE="STATIC" ./scheduling
```

[illegible]


```
OMP_NUM_THREADS=4 OMP_SCHEDULE="STATIC,2" ./scheduling
```

```
OMP_NUM_THREADS=4 OMP_SCHEDULE="STATIC,2" ./scheduling
```

[illegible]

```
OMP_NUM_THREADS=4 OMP_SCHEDULE="DYNAMIC" ./scheduling
```

[illegible]

```
OMP_NUM_THREADS=4 OMP_SCHEDULE="GUIDED" ./scheduling
```

19

How do you do it for Mandelbrot?

```

1 void mandelbrot_draw(int x_resolution, int y_resolution, int max_iter,
2                     double view_x0, double view_x1, double view_y0, double view_y1,
3                     double x_stepsize, double y_stepsize,
4                     int palette_shift, unsigned char (*image)[x_resolution][3],
5                     int num_threads) {
6     // ....
7
8     #pragma omp parallel num_threads(num_threads)
9     {
10         #pragma omp for schedule(dynamic)
11         for (int i = 0; i < y_resolution; i++)
12         {
13             for (int j = 0; j < x_resolution; j++)
14             {
15                 //pixel calculation ...
16             }
17         }
18     }
19 {

```

Assignment 4 - Edge detection (OpenMP)

Assignment 3 - Edge detection (OpenMP)

- You have two weeks time for this assignment
- Use OpenMP
- no valgrind, helgrind, #threads...
- The speedup with 32 cores must be at least 16
- Consider:
 - Previous strategies may apply here

Assignment 4 - Edge detection (OpenMP) - x_gradient()

```
1 template <typename SrcView, typename DstView>
2 void x_gradient(const SrcView &src, const DstView &dst, int num_threads)
3 {
4     typedef typename channel_type<DstView>::type dst_channel_t;
5
6     for (int y = 0; y < src.height(); ++y)
7     {
8         typename SrcView::x_iterator src_it = src.row_begin(y);
9         typename DstView::x_iterator dst_it = dst.row_begin(y);
10
11         for (int x = 1; x < src.width() - 1; ++x)
12         {
13             static_transform(src_it[x - 1], src_it[x + 1], dst_it[x],
14                             halfdiff_cast_channels<dst_channel_t>());
15         }
16     }
17 }
```

Assignment 4 - Edge detection (OpenMP) - Provided Files

- Makefile
 - contains rules to build executables
 - available targets: parallel, sequential, all (default), clean
 - 'mode=debug make [target]' to build debug version, use 'make clean' before
- main.c
 - main function - argument handling + file handling + call `x_luminosity_gradient()`
 - `x_luminosity_gradient()` calls `x_gradient()`
 - you implement the parallel version of `x_gradient()`
- `x_gradient.h`
 - Header file for `x_luminosity_gradient()`
- `x_gradient_seq.h`
 - Sequential version of `x_gradient()`
- `student/x_gradient_par.h`
 - Implement the parallel version in this file
- `unit_test.c`
 - The unit tests that execute both the serial and parallel version to compare results.

Assignment 4 - Edge detection (OpenMP) - Compilation and execution

- Compilation
 - You need to install libjpeg, boost and boost/gil
 - `make [all] [sequential] [parallel] [unit_test]`
 - You implement your solution in a header file
 - You have to make `clean` every time and make again
- Execution
 - `./student/x_gradient_seq`
 - `./student/x_gradient_par -t 4 -f tum.jpg`
 - `./student/unit_test`